



Classical Classification

FinTech
Lesson 11.1



Class Objectives

In today's class we'll learn about classification algorithms



Logistic regression



Support vector machines

Use Cases:



Fraud detection



Price and ROI forecasting



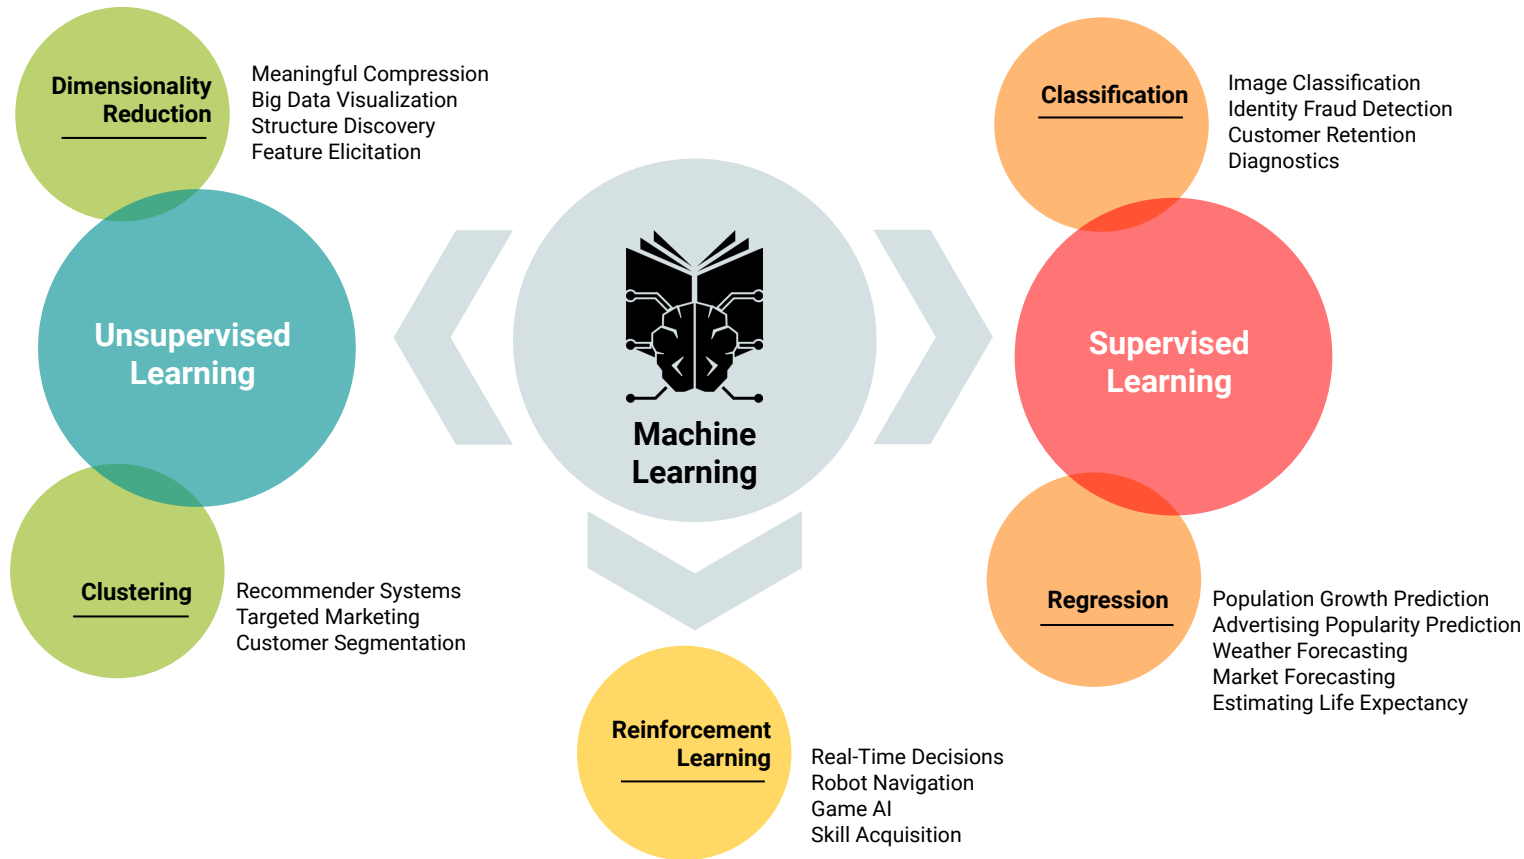
Medical disease/condition diagnosis




Instructor Demonstration


Demo Homework

This is the second week of machine learning!



Intro to Classification






Classification is the action or process of categorizing something according to shared qualities or characteristics.

Classification


Classification is the prediction of discrete outcomes. Outcomes are identified as labels/discrete outputs, which serve to categorize bi-class and multi-class features.




Classification

There are multiple approaches to classification. These include:

Logistic Regression



Support Vector Machines



Classification

Classification is used to forecast and predict financial outcomes, automate underwriting and insurance premiums, detect and categorize health issues and overall health.



Classification

Classification models have drastically improved financial efforts to properly categorize applicants, predict market decline, and categorize fraudulent transactions or suspicious activity.



Classification


FICO credit scoring uses a classification model for its cognitive fraud analytics platform. Classification engines have allowed the financial industry to become more effective and efficient at mitigating risk.



Making Predictions with Logistic Regression


Making Predictions with Logistic Regression

Logistic regression is a common approach used to classify data points and make predictions.




Making Predictions with Logistic Regression

Predictions are made by creating linear paths between data points.




Making Predictions with Logistic Regression

Data points along the trajectory are normalized between 0 and 1.



Making Predictions with Logistic Regression

If a value is above a certain threshold, the data point is considered either of class 0 or 1.



Logistic Regression Model

Running a logistic regression model involves 4 steps, which can be applied when running any machine-learning model:

01

Preprocess

02

Train

03

Validate

04

Predict



Instructor Demonstration

Logistic Regression using Scikit-Learn



Activity: Predicting Diabetes

In this activity, you will use the sklearn library to execute logistic regression models in order to predict whether or not an individual has diabetes.

Suggested Time:
15 minutes





Time's Up! Let's Review.

Review: Predicting Diabetes



How well did your model perform?



How do you know? Did you count the results?



If you were asked to diagnose a patient, how confident would you be in your model's prediction?

Evaluating Logistic Regression Predictions



**How sure are you that
models can actually
predict diabetes?**

Answer

75%

sure, as
described by
the scored
accuracy.






Would you feel comfortable giving the diagnosis of diabetes based off the predictions of the model.



No. The prediction is not 100% accurate. There is room for error, as well as false positives.

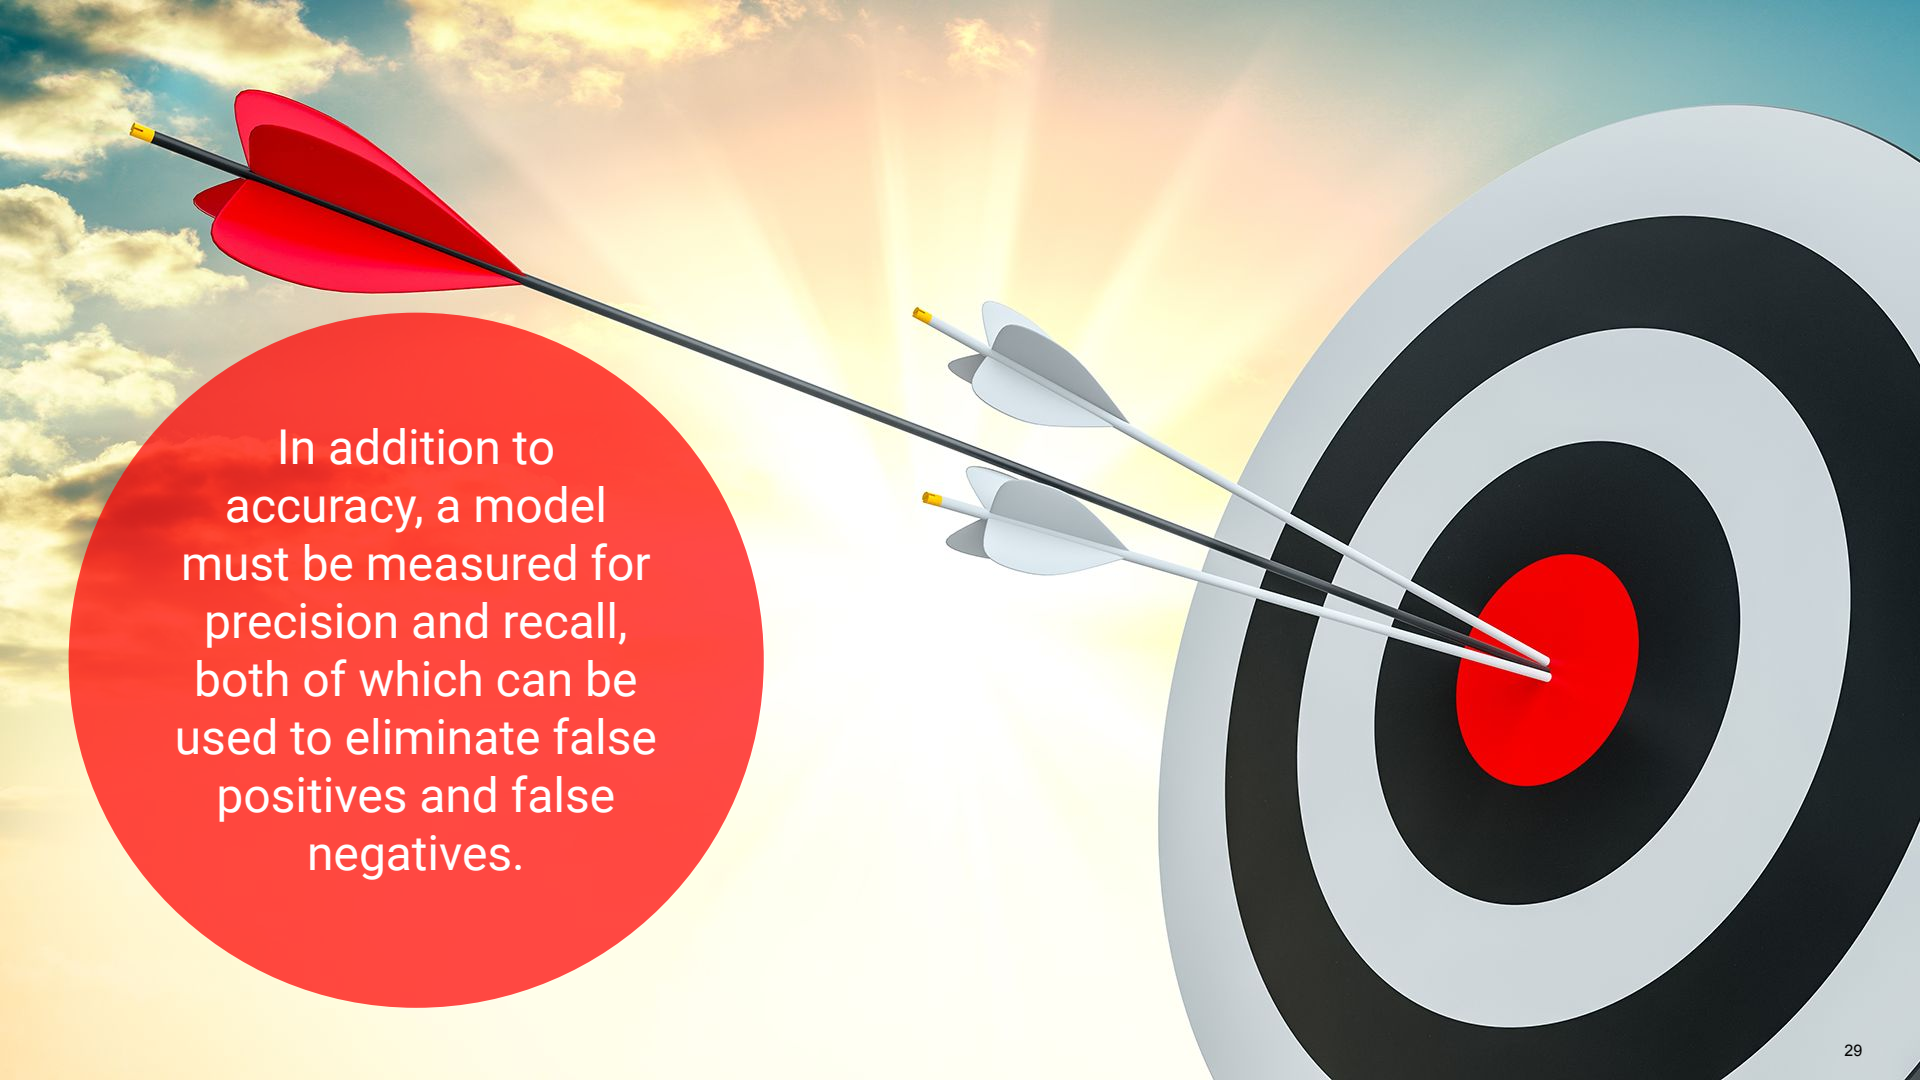


**What is better:
the false positive
or false negative?**

Answer

False Positive. Additional tests can be ran to refine the prediction and filter out individuals who do not have diabetes. This way, those with the potential of having it can be given the treatment and attention they need.




The background of the slide is a vibrant sunset or sunrise scene with a bright sun low on the horizon, casting long, golden rays across a sky filled with soft, white clouds. On the right side of the image, a large target is visible, featuring concentric rings of black, white, and grey, with a prominent red bullseye in the center. Three arrows are shown in flight, all of which have successfully hit the red bullseye. One arrow is red with a black shaft, while the other two are white with black shafts. On the left side, a large red circle contains white text.

In addition to accuracy, a model must be measured for precision and recall, both of which can be used to eliminate false positives and false negatives.

Accuracy, Precision, Recall


Accuracy, Precision, Recall


Accuracy, precision, and recall are especially important for classification model that involve a binary decision problem. Binary decision problems have two possible correct answers: **True Positive** and **True Negative**.



Accuracy, Precision, Recall

Inaccurate and imprecise models result in models returning false positives and false negatives.






Accuracy is how often the model is correct—the ratio of correctly predicted observations to the total number of observations.

Accuracy

Scoring will reveal how accurate the model. However, it does not communicate how precise it is.




Accuracy

Accuracy can be very susceptible to imbalanced classes. In the case of the homework assignment, the number of good loans greatly outweighs the number of at-risk loans. In this case, it can be really easy for the model to only care about the good loans because that has the biggest impact on accuracy. However, we also care about the at-risk loans, so we need a metric that can help us evaluate each class prediction.

Calculation:

$$(TP + TN) / (TP + TN + FP + FN)$$



Precision is the ratio of correctly predicted positive observations to the total predicted positive observations.

Precision

Another example of precision is of all of the individuals that were classified by the model as being a credit risk, how many actually were a credit risk?

The question at hand: Did we classify comprehensively and correctly?




Precision

High precision relates to a low false positive rate.

Calculation:

$$TP / (TP + FP)$$




Recall is the ratio of correctly predicted positive observations to all predicted observations for that class

Recall

Of all of the actual diabetes/credit risk samples, how many were correctly classified as having diabetes/being a credit risk.

The question at hand: Did we classify all samples correctly, leaving little room for false negatives.




Recall

High recall relates to a more comprehensive output and a low false negative rate.

Calculation:

$$TP / (TP + FN)$$

Confusion Matrix & Classification Report



A **confusion matrix** is used to measure and gauge the success of a model.

Confusion Matrix

Confusion matrices reveal the number of true negatives and true positives (actuals) for each categorical class and compares it to the number of predicted values for each class.

n=165	Predicted: No	Predicted: Yes
Actual=No	50	10
Actual=Yes	5	100

Confusion Matrix


These values are then individually summed by column and row. The aggregate sums are then compared to gauge accuracy and precision. If the aggregates match, the model can be considered accurate and precise.

n=165	Predicted: No	Predicted: Yes	
Actual=No	50	10	=60
Actual=Yes	5	100	=105
	=55	=110	

Classification Report

Classification report identifies the **precision**, **recall**, and **accuracy** of a model for each given class.

	precision	recall	f1-score	support
No Diabetes	0.77	0.90	0.83	125
Diabetes	0.72	0.49	0.58	67
accuracy			0.76	192
macro avg	0.74	0.69	0.71	192
weighted avg	0.75	0.76	0.74	192



Instructor Demonstration

Confusion Matrix & Classification Report



Activity: Diagnosing the Model

In this activity, you will return to the model you created to predict diabetes and will use a confusion matrix and classification report to evaluate and diagnose the model.

Suggested Time:
10 minutes





Time's Up! Let's Review.



Activity: Build Loan Approver

In this activity you will apply the machine learning concepts and technical skills learned thus far to create a model for approving loans.

Suggested Time:
15 minutes






Time's Up! Let's Review.

A close-up, high-angle photograph of a computer keyboard. The central focus is a large, white, rectangular key with rounded corners. On this key, there is a dark blue icon of a coffee cup with three wavy lines above it representing steam. Below the icon, the word "Break" is printed in a dark blue, serif font. The key is set against a light-colored, textured keyboard surface. Surrounding the main key are other keys, including one with a double quote symbol to the left and one with a dash/slash symbol to the right, all slightly out of focus.

Break


Support Vector Machines



Support Vector Machines (SVM) is a supervised learning model that can be used for classification and regression analysis. SVM separates classes of data points into multidimensional space.


Linear Classifiers

Linear classifiers attempt to draw a line that separates the data, but which line best separates the groups?




Linear Classifiers

Linear classifiers attempt to draw a line that separates the data, but which line best separates the groups?




Linear Classifiers

Linear classifiers attempt to draw a line that separates the data, but which line best separates the groups?




Linear Classifiers

Linear classifiers attempt to draw a line that separates the data, but which line best separates the groups?




Linear Classifiers

Linear classifiers attempt to draw a line that separates the data, but which line best separates the groups?




Support Vector Machines

The Support Vector Machines (SVM) algorithm finds the optimal hyperplane that separates the data points with the largest margin possible.




Support Vector Machines

The space is segmented by a line or plane that groups data points into their respective classes.




Support Vector Machines

The goal with hyperplanes is to get the margin of the hyperplane equidistant to the data points for all classes.




Support Vector Machines

The data closest to/within the margin of the hyperplane are called support vectors, and they are used to define boundaries of the hyperplane.




Hyperplanes

Hyperplanes can be used clearly delineate classes in multiple dimensions.




Zero tolerance with perfect partition

Hyperplane also supports what is considered zero tolerance with perfect partition, which is a nonlinear hyperplane that will position and orient the hyperplane to correctly classify overlapping or outlying data points.



Zero tolerance with perfect partition

In order to establish zero tolerance with perfect partition, the SVM model may introduce a new **z-axis** dimension for nonlinear hyperplanes.





Instructor Demonstration

SVM model with sklearn

SVM model

Steps to implement an SVM model include:

01

Create the model with appropriate `kernel` parameters

02

Fit the model

03

Extract min and max decision boundaries and store in a mesh grid

04

Execute the `decision_function` to get classifier scores for pre-existing data points

05

Run the `predict` function to classify new data points



Activity: SVM Loan Approver Activity Review

In this activity you will update your loan approver with an SVM model and rerun the evaluation metrics.

Suggested Time:
15 minutes





Time's Up! Let's Review.




Which Model is the Best?


Which Model is the Best?

Both the Logistic Regression and SVM models were both able to predict outcomes; however, the important question is **which model performed best?**

Logistic Regression



Support Vector Machines





**Which is the best approach
to evaluate both models.**

Answer:

Compare the confusion matrices and classification reports.

Confusion Matrices:

n=165	Predicted: No	Predicted: Yes	
Actual= No	50	10	=60
Actual= Yes	5	100	=105
	=55	=110	

Classification Reports:

	precision	recall	f1-score	support
No Diabetes	0.77	0.90	0.83	125
Diabetes	0.72	0.49	0.58	67
accuracy			0.76	192
macro avg	0.74	0.69	0.71	192
weighted avg	0.75	0.76	0.74	192

What is the best approach to evaluate both models?

Logistic Regression Loan Approver Classification Report

	precision	recall	f1-score	support
approve	0.44	0.33	0.38	12
deny	0.50	0.62	0.55	13
micro avg	0.48	0.48	0.48	25
macro avg	0.47	0.47	0.47	25
weighted avg	0.47	0.48	0.47	25

SVM Loan Approver Classification Report

	precision	recall	f1-score	support
approve	0.58	0.58	0.58	12
deny	0.62	0.62	0.62	13
accuracy			0.60	25
macro avg	0.60	0.60	0.60	25
weighted avg	0.60	0.60	0.60	25

What is the best approach to evaluate both models?

The SVM model performed best. **Precision**, **recall**, and **accuracy** were all higher for the SVM loan approver.

	precision	recall	f1-score	support
approve	0.44	0.33	0.38	12
deny	0.50	0.62	0.55	13
micro avg	0.48	0.48	0.48	25
macro avg	0.47	0.47	0.47	25
weighted avg	0.47	0.48	0.47	25



	precision	recall	f1-score	support
approve	0.58	0.58	0.58	12
deny	0.62	0.62	0.62	13
accuracy			0.60	25
macro avg	0.60	0.60	0.60	25
weighted avg	0.60	0.60	0.60	25

What is the best approach to evaluate both models?

Recall percentage for deny is the same for the SVM and logistic regression loan approver, meaning both algorithms correctly predicted the same number of true positive denials.

	precision	recall	f1-score	support
approve	0.44	0.33	0.38	12
deny	0.50	0.62	0.55	13
micro avg	0.48	0.48	0.48	25
macro avg	0.47	0.47	0.47	25
weighted avg	0.47	0.48	0.47	25

	precision	recall	f1-score	support
approve	0.58	0.58	0.58	12
deny	0.62	0.62	0.62	13
accuracy			0.60	25
macro avg	0.60	0.60	0.60	25
weighted avg	0.60	0.60	0.60	25



Questions?